

Assignment 3

Yassine ZAOUI

yassine.zaoui@ensta-paris.fr

001	1. Baseline Model and Initial Results		
002	The baseline model consisted of a simple neural network,		
003	which exhibited heavy underfitting. Both test and validation		
004	accuracies increased marginally during the early epochs and		
005	plateaued at below 10%, even after 20 epochs. This indi-		
006	cated that the architecture was too simplistic to learn mean-		
007	ingful features for the classification task.		
008	2. Improved Architectures		
009	2.1. Adding Convolutional Layers		
010	To enhance feature extraction, we proposed a deeper archi-		
011	tecture with 5-6 convolutional layers. Each layer employed		
012	primarily 3×3 kernels for computational efficiency, with		
013	occasional 5×5 kernels to expand the receptive field. Out-		
014	put channels were doubled after each layer, starting with		
015	32 and ending with 512, closely matching the number of		
016	classes.		
017	2.2. Global Average Pooling and Activation Func-		
018	tions		
019	The final convolutional layer was followed by global av-		
020	erage pooling, reducing the feature dimensionality to 512		
021	before classification. LeakyReLU activation with a nega-		
022	tive slope of 0.001 was used to prevent dead neurons and		
023	maintain gradient flow.		
024	2.3. Results and Challenges		
025	Despite these changes, the model displayed overfitting:		
026	training accuracy exceeded 96%, while validation accuracy		
027	stagnated around 50%, even after 30 epochs.		
028	3. Mitigating Overfitting: Data Augmentation		
029	and Batch Normalization		
030	To improve generalization, we applied data augmentation		
031	techniques, including random rotations, flips, and scaling.		
032	Batch normalization was incorporated after each convolu-		
033	tional layer to stabilize the learning process by normalizing		
034	feature distributions.		
	Improved Results		035
	These strategies significantly improved validation accu-		036
	racy, which reached 63%, while training accuracy remained		037
	around 95%. However, the overfitting issue persisted.		038
	4. Exploration of Pretrained Models		039
	To further enhance performance, pretrained models were		040
	explored. Initially, the <code>DinoV2</code> model was explored for		041
	its state-of-the-art capabilities in feature extraction. While		042
	it effectively avoided overfitting, fine-tuning was com-		043
	putationally expensive and resulted in validation accu-		044
	racy stagnating at approximately 50%-60%. Consequently,		045
	<code>ResNet-50</code> was chosen as an alternative. We took its ar-		046
	chitecture, replaced its fully connected layer with a new one		047
	having 500 outputs. We introduced also a dropout layer		048
	with fairly high rate 0.5 to reduce overfitting.		049
	4.1. Training modifications		050
	We introduced Mixup regularization to encourage		051
	generalization over overfitting, AdamW with bet-		052
	ter weight regularization and fast convergence, Co-		053
	sine Annealing Warm restarts to dynamically adjust		054
	the learning rate, Mixed-Precision training using		055
	<code>torch.cuda.amp.GradScaler</code> , early stopping		056
	to detect when we fail to improve the validation accuracy		057
	and finally label smoothing to the loss function to avoid		058
	being over confident predictions.		059
	4.2. Performance		060
	This approach achieved a better validation accuracy of		061
	roughly 70% proof of improvement. Nevertheless, no one		062
	denies some overfitting persists and we need further exper-		063
	imentation to mitigate this problem.		064
	References		065